

Definitions

What is ML ?


- **a definition:**

For a given **Task T**, a **machine** (algorithm) **A** obtains better **performance P** after an **experiment E**. (It has ***learned*** from it)
(Experiment ~ data)

Yann LeCun's cake metaphor:

- **3 types** of learning :

- **Supervised:** w/ labels
- **Unsupervised:** w/o labels (incl. self-supervised)
- **Reinforcement** (outside this course)



Reinforcement

Supervised

Unsupervised

Today – Outline

- **Supervised Learning basics:**
 - Linear **regression**
 - Polynomial regression
- Lots of **Vocabulary**, notations
- Optimization basics: **Gradient Descent**
- **Supervised Learning**
 - Classification with the Perceptron (maybe)

Today:

Supervised Learning

Input: $\vec{x}^{(n)} = (x_d^{(n)})_{d \in [1, \dots, D]}$, $X = \{\vec{x}^{(n)}\}_{n \in [1, \dots, N]}$

- Expected Output: y^{GT} or $t^{(n)}$ (**Ground Truth**)

Which kind of Task \rightarrow depends on $t^{(n)}$

- **Model:** $y^{predicted} \equiv \hat{y}^{(n)} = \sigma(f_{\Theta}(\vec{x}^{(n)}))$

fct. f_{Θ} is **parameterized** by **parameters**

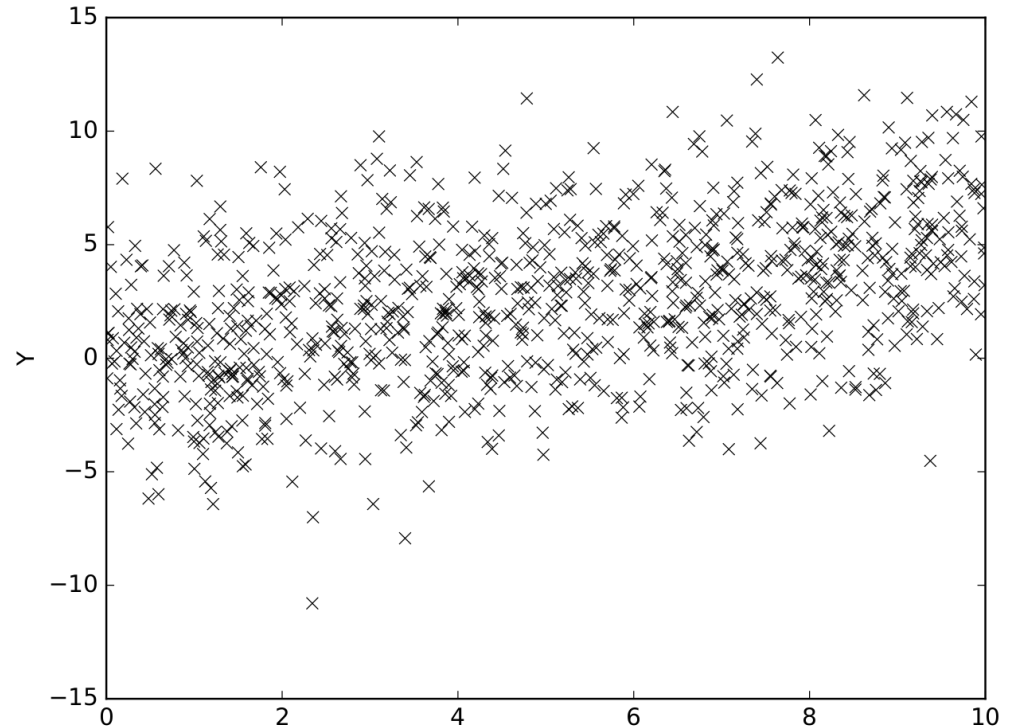
- **Learning** : finding *optimal* parameters to minimize discrepancy between \hat{y} and Ground Truth t

$$\Theta^* = \operatorname{argmin}_{\Theta} \left(\sum_n^N \mathcal{L}(\hat{y}_n, t_n) \right)$$

- **Cost Function (loss function)** : to be chosen ₃

Supervised Learning: Regression

Pairs of data
points $\vec{x}^{(n)} = (x_1^{(n)}, x_2^{(n)})$



→ Relationship $f(x)=y$?

→ **Regression**

- **linear:** $f_{a,b}(x) = ax + b$ or $f_{\vec{a},b}(\vec{x}) = \vec{a} \cdot \vec{x} + b$

- **polynomial:** $f_{\Theta}(\vec{x}) = \vec{\theta} \cdot \Phi(\vec{x})$

(degree P) (see polynomial feature maps)

More Vocabulary

(+case of Regression)

Input: $\vec{x}^{(n)} = (x_d^{(n)})_{d \in [1, \dots, D]}, X = \{\vec{x}^{(n)}\}_{n \in [1, \dots, N]}$

• *Ground Truth:* $t^{(n)} \in \mathbb{R}, T = \{t^{(n)}\}_{n \in [1, \dots, N]}$

Continuous output \rightarrow Task is **Regression**

• **Model:** e.g. a linear function of the input : $f_{\vec{a}, b}(\vec{x}) = \vec{a} \cdot \vec{x} + b$

- Parameters: $\Theta = \{b, a_d; d = 1, \dots, D\}$

- **Prediction:** simply $\hat{y}_n = f_{\Theta}(\vec{x}_n)$



• Learning **Algorithm:**

$$Card(\Theta) = 1 + D$$

- **Initialization:** $\Theta = \Theta_0$

- Minimize some Loss $\mathcal{L}(\hat{y}_n, t_n)$ (to choose)

- For this, use some minimization scheme (Grad. Desc.)

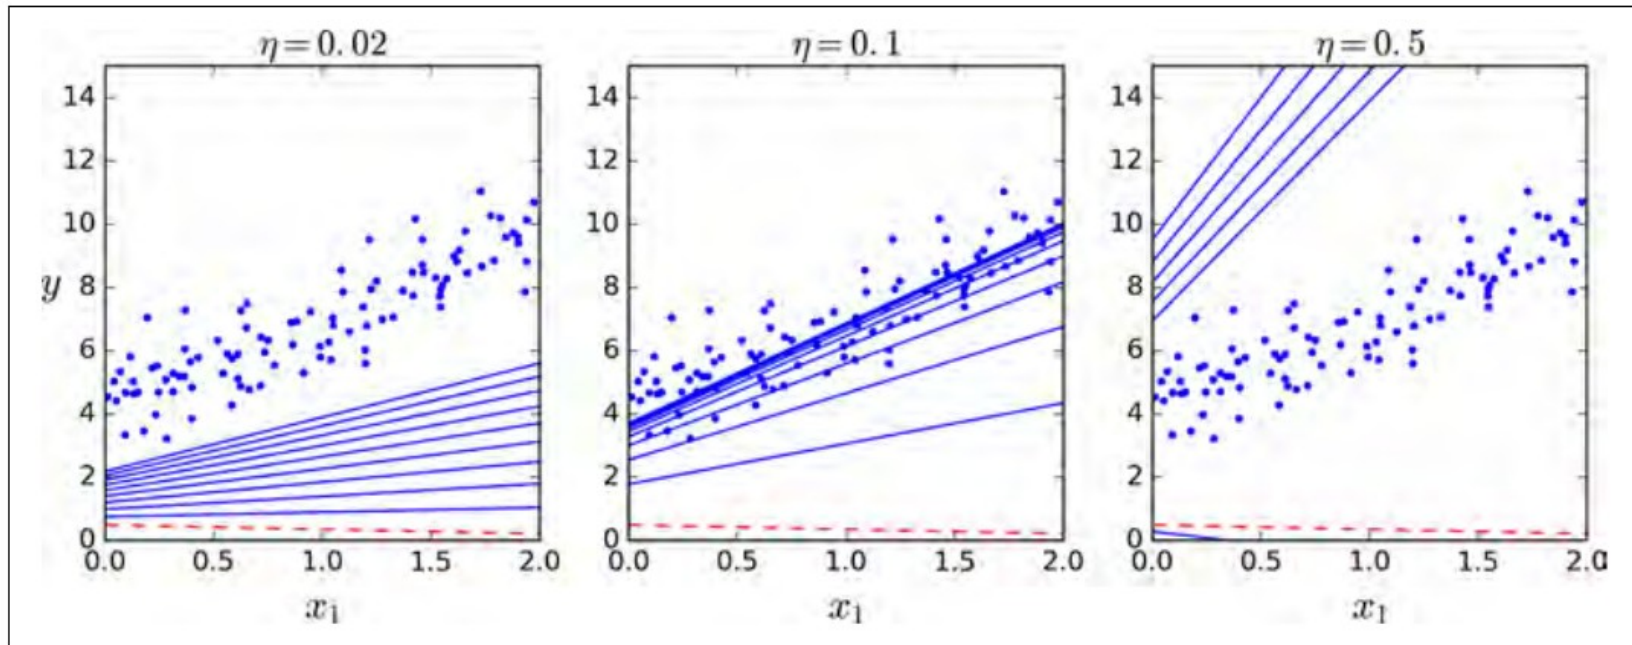
Supervised Learning: Regression

- We can choose: **Least Squares**

Single data point Loss: $\mathcal{L}(f_{\Theta}(\vec{x}_n), t_n) = (\vec{f}(\vec{x}^{(n)}) - t^{(n)})^2$

$$\text{Global Loss: } \mathcal{L}(X, T) = \frac{1}{N} \sum_{n=1}^N \mathcal{L}(f_{\Theta}(\vec{x}_n), t_n)$$

- **Gradient Descent:**









Gradient Descent


short reminder

- I have a function $J(\theta)$ and want to find the value θ^* for which $J(\theta)$ is minimum


What is the gradient ?




$x^3 + y^3$, et $-(x^2 + y^2) + (x^4 + y^4)$:



$x^2 + y^2$ et $||\vec{x}|| - a\vec{w} \cdot \vec{x} = (x^2 + y^2)^{1/2} - aw_1x - aw_2y$, avec $a = 3, w_1 = 0.1, w_2 = 0.3$:



$e^{-x^2}y^2$, et $-(x^2 + y^2) + (x^2 + y^2)^2$



Gradient Descent

- It goes in the steepest direction (from the local point) → is also called “*steepest descent*”
- Limitations:
 - at best, converges to ***one of the local minima***
 - *typically* converges to the **local attractor** (min. in the local basin of attraction)
 - Result **depends on starting position** !
 - it **may never converge** ! (diverge or continuously go down)

Least Squares

$$LSE = \frac{1}{N} \sum_{n=1}^N (\vec{f}_{\Theta}(x^{(n)}) - \vec{y}^{(n)})^2 \quad , \text{ with } f_{\Theta}(\vec{x}) = \vec{\theta} \cdot \vec{x}$$

Trick: Augmented data

- Add 1's into X to take care of the offset, once and for
→ get cleaner equations (and cleaner code) !

A word on unsupervised: the example of K-means

- Goal: find groups in data (X). No labels (y).
- Idea: assign “classes” (assignment to a *group*, aka *cluster*) to points anyway, and ask to have **homogeneous groups**:
 - close-by points should belong to the same cluster
 - clusters should be batches of points which are close enough
- In practice: cook a cost function J that realizes this, then minimize it.

$J =$

- Numerical minimization is performed approximately, by starting at random, then iterating 2 steps:
 - each point is assigned to the closest cluster center
 - each cluster center is the barycenter of the data points assigned to it

References:

Linear regression (by G.D.)

→ Bishop book, page 143-144, section 3.1.3
(sequential learning)

→ https://en.wikipedia.org/wiki/Least_squares#Linear_least_squares

- **Gradient Descent (assumed known)**

→ catch up lesson:

https://en.wikipedia.org/wiki/Gradient_descent

Key concepts

- **Supervised Learning**
- **Regression**
- **Task, Model, parameters, prediction/decision, input feature**